

DSA – MODULE-1

TextHack as a Text Analytics and Sequence Search Engine

- ❖ TextHack is a **Text Analytics** and **Sequence Search Engine** that provides efficient algorithms for searching, analyzing, and processing textual and sequence data. It is an educational platform that demonstrates how different algorithmic techniques solve real-world text processing problems.
- ❖ Instead of focusing on a single algorithm, TextHack integrates several algorithms to perform tasks such as **pattern searching**, **fuzzy matching**, **document similarity analysis**, **citation-flow analysis**, **scheduling demonstrations**, and **large-prime testing**.
- ❖ It is widely used for learning **string algorithms**, **dynamic programming**, **graph algorithms**, **approximation algorithms**, and **randomized algorithms**.

Features of TextHack

1. Pattern Search

2. Fuzzy Matching

3. Document Similarity Analysis

4. Citation-Flow Analysis

5. Scheduling Demonstrations

6. Large-Prime Testing

1. Pattern Search

Definition

Pattern search is the process of **finding occurrences** of a **specific word, phrase, or sequence** within a large collection of text or biological sequences.

TextHack supports both **exact matching** and **efficient sequence searching**.

Example

Suppose the document contains:

```
Computer Science is the study of algorithms.  
Algorithms solve computational problems.
```

Search Pattern:

```
Algorithms
```

Output:

```
Match 1: Line 1
```

```
Match 2: Line 2
```

Algorithms Used

- **Naïve** String Matching
- Knuth-Morris-Pratt (**KMP**)
- **Rabin-Karp**
- Boyer-Moore
- Suffix Trees/Arrays

Applications

- Search engines
- DNA sequence analysis
- Digital libraries
- Log analysis
- Source code search

1. Pattern Search

Naive search compares every position.

Time complexity: $O(n \times m)$.

Efficient algorithms: KMP, Z-Algorithm, Rabin–Karp.

Observation: Search engines cannot afford to scan character-by-character for every query.

2. Fuzzy Matching

Definition

Fuzzy matching finds **approximately matching strings** rather than requiring an **exact match**. It **tolerates spelling mistakes**, missing characters, and typographical errors.

Example

Database:

Hyderabad

Algorithm

Computer

User searches:

Algoritm

Output:

Algorithm

Similarity = 95%

Common Algorithms

- **Levenshtein Distance**
- Hamming Distance
- Jaro-Winkler Similarity
- Soundex

Applications

- Spell checking
- Duplicate record detection
- OCR correction
- Customer database matching
- Search engines

2. Fuzzy Matching

User types: “algoritm” instead of “algorithm”.

How many edits are needed?

Notice that algoritm is missing the letter **h**.

algoritm → algorithm = 1 insertion

Dynamic Programming

3. Document Similarity Analysis

Definition

Document Similarity is the process of measuring **how similar two or more documents are based on their content**. The result is usually expressed as a **similarity score** between **0 and 1** (or **0% to 100%**).

1 (100%) → Documents are identical.

0 (0%) → Documents are completely different.

Example

Document A

Machine learning improves healthcare.

Document B

Machine learning helps healthcare.

Similarity

Jaccard Similarity = $3/4$

Techniques

- Bag of Words
- TF-IDF
- Cosine Similarity
- Jaccard Similarity
- Semantic Embeddings
(BERT, Word2Vec)

Applications

- Plagiarism detection
- Document clustering
- Search engines
- Recommendation systems
- Duplicate document detection

3. Document Similarity Analysis

DocA “Machine learning improves predictions”

DocB “Machine learning improves forecasting”

So how to compute similarity?

3. Document Similarity Analysis

Idea: Convert documents into sets of words and compute similarity.

Jaccard Similarity:

Example: Common words = {machine, learning, improves}

$$\text{Similarity} = 3 / 5 = 0.6$$

Use: Plagiarism detection, recommendation systems.

4. Citation-Flow Analysis

Definition

Citation-flow analysis studies how **knowledge and ideas spread through citations** among research papers.

Documents are represented as a directed graph.

- Nodes $\rightarrow$ Research papers
- Edges $\rightarrow$ Citations

Example

Paper A



Paper B



Paper C

Paper B cites Paper A.
Paper C cites Paper B.
Knowledge flows from $A \rightarrow B \rightarrow C$.

Applications

- Identify influential papers
- Research trend analysis
- Author impact measurement
- Academic recommendation systems


Advantages

- Tracks knowledge evolution
- Finds highly cited publications
- Supports bibliometric analysis

4. Citation-Flow Analysis

In **TextHack**, citation flow analysis is one of the advanced functionalities. It allows users to:

- Find influential research papers.
- Identify the most cited authors.
- Trace the evolution of a research topic.
- Discover connections between different research areas.
- Recommend relevant papers based on citation relationships.



Represent papers as a graph:

Vertex: Research paper

Edge: Citation

Questions:

Which paper is most influential?


How does information flow?

Which papers form a research community?

Algorithms: BFS, DFS, PageRank-like methods.

4. Citation-Flow Analysis


Consider the following citation graph:



- Which paper has the highest **in-degree**?
- Which paper is likely the **foundational** paper?
- If you start a **BFS** from A, what is the traversal order?
- Which paper combines ideas from **B** and **C**?

4. Citation-Flow Analysis

Consider the following citation graph:



Which paper has the highest **in-degree**?

D (it is cited by both B and C).

Which paper is likely the **foundational** paper?

A (it is the starting point of the citation chain).

If you start a **BFS** from A, what is the traversal order?

A → B → C → D → E.

Which paper combines ideas from **B** and **C**?

D, because it cites or builds upon both B and C.

5. Scheduling Demonstrations

Definition

TextHack includes demonstrations of scheduling algorithms to show how tasks are executed efficiently using different scheduling strategies.

Example

Task	Execution Time
T1	4
T2	2
T3	5

Using **Shortest Job First (SJF)**:

Execution order:

T2 → T1 → T3

Other scheduling algorithms demonstrated:

- First Come First Serve (FCFS)
- Shortest Job First (SJF)
- Priority Scheduling
- Round Robin
- Earliest Deadline First (EDF)

Applications

- Operating systems
- CPU scheduling
- Cloud computing
- Project scheduling

5. Scheduling Demonstrations

Scenario: TextHack must process 5 uploaded documents.

Task	Time
D1	2 min
D2	1 min
D3	5 min

Question: What order minimizes average waiting time?

Greedy Choice: Shortest Job First (SJF)

Order: D2 $\rightarrow$ D1 $\rightarrow$ D3

Application: Cloud servers, print queues, CPU scheduling.

6. Large-Prime Testing (Primality Testing)

Definition

Large-prime testing determines whether a very large number is **prime or composite**.

- Large prime numbers are essential in **secure indexing, hashing, and cryptography**. Some hashing techniques use prime numbers to reduce collisions. This functionality is important in **cryptography**, especially in RSA encryption.

Example

Input

999983

Output

Prime Number

Algorithms Used

- Trial Division (small numbers)
- Fermat Primality Test
- Miller-Rabin Test
- AKS Primality Test

Applications

- RSA key generation
- Digital signatures
- Secure communication
- Cryptographic protocols

6. Large-Prime Testing (Primality Testing)


Need: Secure indexing and hashing.

Question: Is 2,147,483,647 prime?

Naive algorithm → A deterministic algorithm that checks divisibility up to $\sqrt{n}$.

Probabilistic algorithm → A class of algorithms that use randomness and provide very high confidence.

Miller–Rabin → A specific probabilistic algorithm used to test whether a number is prime quickly.



Evaluation of Different Categories of User Queries and the Algorithms Used for Efficient Processing

- ❖ Different types of user queries require different algorithmic techniques depending on the nature of the problem.
- ❖ Search engines, recommendation systems, route planners, scheduling systems, and optimization tools use **String Algorithms, Dynamic Programming, Network Flow Algorithms, Approximation Algorithms, and Randomized Algorithms** to provide efficient and accurate results.

1. String Algorithms

Purpose

String algorithms are used to **search, compare, and manipulate text or sequences efficiently.**

Role in TextHack - User Queries

In TextHack, string algorithms are used to:

- **Search for keywords in documents.**
- **Perform exact pattern matching.**
- **Search DNA or protein sequences.**
- **Index and retrieve text efficiently.**
- **Spell checking and autocomplete.**

Example

User Query

Search "algorithm" in a textbook.

Document

Data Structures and Algorithms are important.

Algorithms improve efficiency.

Output

Match found in Line 1 and Line 2.

Algorithms Used

- Naïve String Matching
- Knuth-Morris-Pratt (KMP)
- Rabin-Karp
- Boyer-Moore

Why They Are Efficient

- Avoid unnecessary comparisons.
- Search large text quickly.
- Support exact and multiple pattern matching.

Applications

- Search engines
- Text editors
- DNA sequence search
- Plagiarism detection

2. Dynamic Programming (DP)

Purpose

Dynamic Programming solves **optimization problems** by **storing solutions to overlapping subproblems**, avoiding repeated computations.

User Queries

- Find the shortest route.
- Find the minimum cost.
- Schedule jobs efficiently.
- Compute edit distance between strings.

Example

User Query

What is the minimum cost to travel from A to D?

Possible paths

Output

Minimum Cost = 6

Path	Cost
A→B→D	8
A→C→D	6

Algorithms Used

- 0/1 Knapsack
- Matrix Chain Multiplication
- Longest Common Subsequence (LCS)
- Edit Distance
- Optimal Binary Search Tree (OBST)

Why It Is Efficient

- Stores intermediate results.
- Eliminates repeated calculations.
- Produces optimal solutions.

Applications

- Route optimization
- Resource allocation
- Sequence alignment
- Speech recognition

2. Dynamic Programming (DP)

Role in TextHack

TextHack uses Dynamic Programming for:

- **Computing edit distance between documents.**
- **Sequence alignment.**
- **Finding longest common subsequences.**
- **Improving document similarity analysis.**

Example

Document 1:

`algorithm`

Document 2:

`algorithms`

DP computes the minimum number of edits required to transform one word into the other.

Output:

Edit Distance = 1

Algorithms Used

- Edit Distance
- Longest Common Subsequence (LCS)
- Longest Common Substring

Applications

- Spell checking
- Document comparison
- DNA sequence alignment
- Text correction


3. Network Flow Algorithms

Purpose

Network Flow algorithms determine the **maximum possible flow** through a network while satisfying capacity constraints.

User Queries

- Allocate jobs to workers.
- Find the maximum network bandwidth.
- Assign students to projects.
- Optimize traffic flow.



Algorithms Used

- Ford-Fulkerson
- Edmonds-Karp
- Dinic's Algorithm

Why They Are Efficient

- Maximize resource utilization.
- Detect bottlenecks.
- Solve assignment and transportation problems.

3. Network Flow Algorithms

Role in TextHack

Although TextHack is mainly a text analytics system, network flow concepts are useful for:

- **Citation-flow analysis.**
- **Resource allocation during document processing.**
- **Data routing in distributed search systems.**
- **Task assignment among processing nodes.**

Applications

- | | |
|-------------------------|---------------------------|
| ❖ Citation analysis | ❖ Internet routing |
| ❖ Network optimization | ❖ Airline scheduling |
| ❖ Distributed computing | ❖ Supply chain management |
| ❖ Resource scheduling | ❖ Bipartite matching |
| | ❖ Traffic management |

4. Approximation Algorithms

Purpose

Approximation algorithms provide **near-optimal solutions** for problems where finding the exact solution is computationally infeasible (NP-hard problems).

User Queries

- Plan delivery routes.
- Select minimum service centers.
- Optimize delivery coverage.
- Solve large scheduling problems.

Example

Travelling Salesman Problem (TSP)

Cities

A B C D

Optimal distance

80 km

Approximation algorithm returns

84 km

The solution is close to optimal and computed much faster.

Algorithms Used

- Vertex Cover Approximation
- Set Cover Greedy Algorithm
- Christofides Algorithm (TSP)
- Greedy Scheduling Algorithms

Why They Are Efficient

- Fast execution.
- Scalable to large inputs.
- Guaranteed performance bounds.

4. Approximation Algorithms

Role in TextHack

TextHack uses approximation algorithms for:

- **Large document clustering.**
- **Search result ranking.**
- **Topic grouping.**
- **Efficient scheduling of analysis tasks.**

Example

Suppose there are millions of documents. Instead of finding the perfect grouping, TextHack quickly produces a grouping that is very close to the optimal one.

Applications

- Large-scale search
- Document clustering
- Recommendation systems
- Resource optimization
- Logistics
- Vehicle routing
- Cloud computing
- Facility location
- Network design

5. Randomized Algorithms

Purpose

Randomized algorithms use **random choices** during execution to improve average performance or simplify computation.

User Queries

- Generate secure cryptographic keys.
- Random sampling.
- Test whether a large number is prime.
- Shuffle data randomly.

Example

QuickSort

Input

8 5 2 9 4 7

Randomly selected pivot

2

The algorithm partitions the array and sorts efficiently.

Algorithms Used

- Randomized QuickSort
- Miller-Rabin Primality Test
- Randomized Hashing
- Monte Carlo Algorithms
- Las Vegas Algorithms

Why They Are Efficient

- Good average-case performance.
- Reduces worst-case behavior.
- Handles large datasets efficiently.

5. Randomized Algorithms

Role in TextHack

TextHack uses randomized algorithms for:

- **Large-prime testing.**
- **Random sampling of documents.**
- **Fast indexing.**
- **Efficient hash-based searching.**

Example

Randomized QuickSort

Input

8 5 3 9 2

Randomly choose pivot

5

The array is partitioned and sorted efficiently.

Algorithms Used

- Randomized QuickSort
- Miller-Rabin Primality Test
- Universal Hashing
- Monte Carlo Algorithms

Applications

- Cryptography
- Fast searching
- Database indexing
- Big-data analytics
- Machine learning
- Random sampling
- Gaming and simulations

Large-Prime Testing (Primality Testing)

Basic Method: Trial Division Algorithm

- 1.If $n \leq 1$, it is not prime.
- 2.Check divisibility from **2** to $\sqrt{n}$.
- 3.If any number divides n exactly, it is **Composite**.
- 4.Otherwise, it is **Prime**.

Example 2: Is 91 Prime?

Step 1

Number

91

Step 2

$$\sqrt{91} \approx 9.54$$

Check divisors:

2–9

Step 3

Divisor	91 ÷ Divisor	Result
2	45.5	Not divisible
3	30.3	Not divisible
4	22.75	Not divisible
5	18.2	Not divisible
6	15.17	Not divisible
7	13	✓ Divisible

Since

$$91 = 7 \times 13$$

Result

91 is Composite.

Comparison of Algorithms

Algorithm Category	Purpose	Role in TextHack	Typical Algorithms	Example Applications
String Algorithms	Search and pattern matching	Find keywords and sequences efficiently	Naïve, KMP, Rabin-Karp, Boyer-Moore	Search engines, DNA search, text retrieval
Dynamic Programming	Find optimal solutions	Edit distance, sequence alignment, document similarity	Knapsack, LCS, Edit Distance, OBST	Spell checking, plagiarism detection
Network Flow Algorithms	Optimize flow in networks	Citation-flow analysis, resource allocation, distributed processing	Ford-Fulkerson, Edmonds-Karp, Dinic	Citation networks, task scheduling
Approximation Algorithms	Near-optimal solutions for NP-hard problems	Document clustering, ranking, large-scale optimization	Vertex Cover, Set Cover, Christofides	Recommendation systems, clustering
Randomized Algorithms	Use randomness for efficient computation	Prime testing, indexing, random sampling, fast sorting	Randomized QuickSort, Miller-Rabin	Cryptography, search optimization, big-data processing

Mapping Queries to Algorithms

“If a user asks ‘Find all papers similar to this one’, which algorithm category should we use?”

Mapping Queries to Algorithms

“If a user asks ‘Find all papers similar to this one’, which algorithm category should we use?”

Answer: Dynamic Programming / Similarity Algorithms

Query to Algorithm Mapping

Exact Search → String Algorithms

Spell Check → DP

Citation → Graph

Scheduling / Task Processing →

Greedy

Prime Test → Randomized

Summary

TextHack integrates multiple algorithm families.

- Exact search → String algorithms.
- Spelling correction → Dynamic programming.
- Citation analysis → Graph algorithms.
- Task processing → Greedy/Approximation.
- Prime testing → Randomized algorithms.

Quiz

Q1. Which algorithm category is most suitable for exact text search?

String Algorithms

Sorting

Searching Trees

Greedy

Quiz

Q1. Which algorithm category is most suitable for exact text search?

String Algorithms

Sorting

Searching Trees

Greedy

Quiz

Q2. Which problem is solved using Edit Distance?

Prime testing

Spelling correction

Shortest path

Scheduling

Quiz

Q2. Which problem is solved using Edit Distance?

Prime testing

Spelling correction

Shortest path

Scheduling

Quiz

Q3. Citations between papers are naturally modeled as a:

- Array
- Stack
- Graph
- Heap

Quiz

Q3. Citations between papers are naturally modeled as a:

Array

Stack

Graph

Heap

Analysis of the Limitations of Linear Search, Conventional Graph Traversal and Basic Optimization Approaches

- a) Basic algorithms such as **Linear Search, Conventional Graph Traversal (BFS and DFS)** and **Basic Optimization Approaches** work well for small datasets and simple problems.
- b) However, as data size and problem complexity increase, these methods become inefficient in terms of execution time and memory usage.
- c) In such situations, advanced algorithms are required to process large-scale datasets efficiently and produce optimal or near-optimal solutions.

1. Limitations of Linear Search

What is Linear Search?

Linear search examines each element one by one until the desired item is found or the entire list has been checked.

Limitations

1. Slow for Large Datasets

If a dataset contains millions of records, linear search may need to examine almost every element.

Example

Searching for a student record among **10 million records**.

Worst case:

- Compare all 10 million records.

3. Not Suitable for Real-Time Systems

Applications like:

- Google Search
- Online shopping websites
- Banking systems

Example

Array:

15 24 38 42 50 63 75

Search for **63**

15 → 24 → 38 → 42 → 50 → 63

The algorithm compares each element sequentially.

2. Time Complexity

Execution time increases linearly with the number of elements.

Case	Complexity
Best Case	$O(1)$
Average Case	$O(n)$
Worst Case	$O(n)$

Better Alternatives

- Binary Search (sorted data)
- Hash Tables
- Balanced Binary Search Trees
- Trie
- B-Tree

2. Limitations of Conventional Graph Traversal (BFS and DFS)

What are BFS and DFS?

- **Breadth-First Search (BFS):** Explores nodes level by level.
- **Depth-First Search (DFS):** Explores one path completely before backtracking.

Limitations

1. Cannot Handle Weighted Shortest Paths Efficiently

Suppose:

A --2--> B

A --8--> C

B --1--> C

BFS chooses the path with the fewest edges:

A → C

Cost = 8

The optimal path is:

A → B → C

Cost = 3

Therefore, BFS is unsuitable for weighted shortest-path problems.

2. Inefficient for Huge Graphs

Modern social networks contain:

- Millions of users
- Billions of connections

Traversing the entire graph becomes expensive.

3. No Capacity Optimization

BFS and DFS cannot solve:

- Maximum flow
- Resource allocation
- Matching problems

Better Alternatives

- Dijkstra's Algorithm
- Bellman-Ford Algorithm
- Floyd-Warshall Algorithm
- Ford-Fulkerson
- Edmonds-Karp

3. Limitations of Basic Optimization Approaches

What are Basic Optimization Approaches?

Basic optimization techniques often rely on simple rules such as:

- Trying every possible solution (Brute Force)
- Greedy decisions
- Exhaustive search

Example

Travelling Salesman Problem (TSP)

Cities:

A B C D E

Brute force checks every possible tour.

For **5 cities**:

$5! = 120$ tours

For **20 cities**:

$20! \approx 2.43 \times 10^{18}$ tours

Impossible to evaluate within a reasonable time.

Limitations

1. Exponential Running Time

Many optimization problems require:

$O(2^n)$ or $O(n!)$ time.

Better Alternatives

- Dynamic Programming
- Branch and Bound
- Approximation Algorithms
- Heuristic Algorithms
- Metaheuristic Algorithms

Situations Where Advanced Algorithms Become Necessary

1. Large-Scale Search Engines

Problem

Search billions of web pages.
Linear search is too slow.

Advanced Algorithms

- KMP
- Boyer-Moore
- Trie
- Inverted Index

Benefit

Search results returned within milliseconds.

2. GPS Navigation

Problem

Find the shortest route among millions of roads.
BFS cannot optimize weighted distances.

Advanced Algorithms

- Dijkstra
- A*
- Floyd-Warshall

Benefit

Fast and accurate route planning.

Situations Where Advanced Algorithms Become Necessary

3. Social Networks

Problem

Analyze billions of user relationships.

DFS/BFS alone are insufficient for ranking or recommendation.

Advanced Algorithms

- PageRank
- Community Detection
- Graph Neural Networks

Benefit

Friend recommendations and influence analysis.

4. Cloud Computing

Problem

Assign thousands of tasks to hundreds of servers.

Simple scheduling causes resource imbalance.

Advanced Algorithms

- Network Flow
- Dynamic Programming
- Approximation Algorithms

Benefit

Improved resource utilization and reduced execution time.

Situations Where Advanced Algorithms Become Necessary

5. E-Commerce Recommendation Systems

Problem

Recommend products from millions of items.

Basic comparison methods become too slow.

Advanced Algorithms

- Machine Learning
- Locality Sensitive Hashing (LSH)
- Approximate Nearest Neighbor Search

Benefit

Personalized recommendations in real time.

6. Cryptography Problem

Test whether a 2048-bit number is prime.

Trial division is impractical.

Advanced Algorithms

- Miller–Rabin Primality Test
- AKS Primality Test

Benefit

Fast generation of secure cryptographic keys.

Basic Approach	Limitations	Advanced Algorithm	Applications
Linear Search	$O(n)$, slow for large datasets	Binary Search, Hash Tables, Tries	Search engines, databases
BFS/DFS	Cannot optimize weighted paths or resource allocation	Dijkstra, A*, Ford-Fulkerson	Navigation, network routing, logistics
Brute Force Optimization	Exponential time, poor scalability	Dynamic Programming, Branch and Bound, Approximation Algorithms	Scheduling, route optimization, resource allocation
Trial Division for Prime Testing	Very slow for large numbers	Miller–Rabin, AKS	Cryptography
Simple Sorting/Searching	Inefficient on massive datasets	Randomized QuickSort, External Sorting	Big data processing

Suitability of Algorithmic Paradigms for TextHack Functionalities

TextHack Functionality	String Algorithms	Dynamic Programming	Network Flow	Approximation Algorithms	Randomized Algorithms	Best Choice
Pattern Search	✓ Excellent	✗	✗	✗	✗	String Algorithms
Fuzzy Matching	✓ (initial search)	✓ Excellent	✗	✗	✗	Dynamic Programming (Edit Distance)
Document Similarity	✓	✓ Excellent	✗	✓ (large datasets)	✗	Dynamic Programming + String Algorithms
Citation-Flow Analysis	✗	✗	✓ Excellent	✓	✗	Network Flow Algorithms
Scheduling Demonstrations	✗	✓	✓	✓ Excellent	✗	Approximation Algorithms
Large-Prime Testing	✗	✗	✗	✗	✓ Excellent	Randomized Algorithms

Introduction to Asymptotic Complexity Analysis

Introduction

When designing algorithms, it is important to know **how efficiently an algorithm performs as the input size increases**. Two major factors determine an algorithm's efficiency:

- 1. Execution Time (Time Complexity)** – How long the algorithm takes to run.
- 2. Memory Requirement (Space Complexity)** – How much memory the algorithm uses.

Asymptotic Complexity Analysis is a mathematical method used to evaluate an algorithm's performance for **large input sizes** without depending on the computer's hardware or programming language. Instead of measuring actual execution time in seconds, asymptotic analysis studies **how the running time and memory grow as the input size (n) increases**.

Types of Asymptotic Notations

1. Big-O Notation (O)

Upper bound: worst-case growth rate. "This algorithm will never do worse than this."

2. Big-Omega Notation (Ω)

Lower bound: best-case growth rate.

3. Big-Theta Notation (Θ)

Tight bound: growth rate when best and worst case coincide (or we care about the typical case).

Complexity	Name	Example Algorithm
$O(1)$	Constant	Array indexing, Stack push/pop
$O(\log n)$	Logarithmic	Binary Search
$O(n)$	Linear	Linear Search
$O(n \log n)$	Linearithmic	Merge Sort, Heap Sort
$O(n^2)$	Quadratic	Bubble Sort, Selection Sort
$O(2^n)$	Exponential	Recursive Fibonacci
$O(n!)$	Factorial	Brute-force Travelling Salesman Problem

Common complexity classes, from best to worst:

$$O(1) < O(\log n) < O(n) < O(n \log n) < O(n^2) < O(n^3) < O(2^n) < O(n!)$$

Execution Time vs. Memory Requirements

Algorithm selection is a constant balancing act between two distinct resources: **Time Complexity** (CPU cycles) and **Space Complexity** (RAM/memory usage). How these two factors influence your choice depends entirely on your system's constraints.

Time Complexity (Execution Time)

This is usually the highest-priority metric for user-facing applications. If an algorithm takes too long, users abandon the app.

- **The Goal:** Minimize operations.
- **The Trade-off:** To make an algorithm faster, you often have to pre-compute data or store extra records, which eats up memory.

Space Complexity (Memory Requirements)

This measures the extra memory an algorithm needs to execute, beyond the space taken up by the input data itself.

- **The Goal:** Prevent memory overflow (like a Stack Overflow Error or running out of RAM).
- **The Trade-off:** Memory-efficient algorithms often have to re-calculate data on the fly, which eats up CPU time.

The Trade-Off in Action

Consider sorting a massive array of numbers. Two common choices highlight this clash:

1. **Merge Sort:** Extremely fast ($O(n \log n)$ time) but requires copying the data into a temporary array ($O(n)$ extra space).
2. **In-Place Quick Sort:** Also fast on average ($O(n \log n)$ time) but sorts the elements right inside the original array, requiring minimal extra memory ($O(\log n)$ stack space).

If you are writing software for a massive cloud server with terabytes of RAM, you will choose Merge Sort because memory is cheap and speed is king. If you are writing software for a tiny medical pacemaker or an IoT smart-plug, memory is highly restricted; you will sacrifice raw speed to save every byte of space.

The Real-World Impact

- ❖ If your dataset is small ($n = 10$), the difference between a Linear algorithm and a quadratic algorithm is negligible (10 vs 100 operations). A computer will execute both instantly.
- ❖ However, if you try to run a quadratic $O(n^2)$ algorithm—like a nested loop checking every item against every other item—on a dataset of 1 million users, your program will attempt to execute **1 trillion operations**. Even on a fast CPU, this could take hours or days, causing the application to hang or crash.
- ❖ In large-scale systems (like Google Search, Netflix recommendations, or financial high-frequency trading), an algorithm that is $O(n^2)$ is fundamentally broken. Scalability forces engineers to relentlessly optimize toward $O(n \log n)$, $O(n)$, or even $O(\log n)$ bounds to ensure that as the business grows, the infrastructure bills and latency don't grow exponentially along with it.

Significance of Scalability in Large Datasets

What is Scalability?

Scalability is the ability of an algorithm to maintain good performance as the size of the input increases.

A scalable algorithm continues to work efficiently for thousands, millions, or even billions of records.

Example

Suppose a search engine stores:

- ☐ 100 web pages
- ☐ 10,000 web pages
- ☐ 10 million web pages
- ☐ 100 billion web pages

A linear search becomes impractical as the dataset grows.

Instead, scalable methods such as:

- ☐ Binary Search
- ☐ Hash Tables
- ☐ Tries
- ☐ Inverted Indexes are used.

Real-World Importance of Scalability

1. Search Engines

Problem:

Search billions of web pages.

Need: Fast search within milliseconds.

Algorithms:

Inverted Index

Trie

String Matching Algorithms

2. Social Networks

Problem:

Analyze billions of user connections.

Need:

Efficient graph processing.

Algorithms:

BFS

DFS

PageRank

Graph Algorithms

3. E-Commerce

Problem:

Search millions of products.

Need:

Fast recommendations.

Algorithms:

Hashing

Dynamic Programming

Approximate Nearest Neighbor Search

4. Cloud Computing

Problem:

Search millions of products.

Need:

Fast recommendations.

Algorithms:

Hashing

Dynamic Programming

Approximate Nearest Neighbor Search

5. Bioinformatics

Problem:

Compare millions of DNA sequences.

Need:

Fast sequence matching.

Algorithms:

KMP

Rabin-Karp

Dynamic Programming

Analysis of Computational Requirements of Search, Similarity Analysis, Citation Analysis, Scheduling, and Primality Testing Functionalities

Introduction

Modern systems such as **TextHack** perform various computational tasks, including **search, similarity analysis, citation analysis, scheduling, and primality testing**. Each functionality has different computational requirements and therefore requires different algorithmic approaches. By analyzing **time complexity** and **space complexity** (asymptotic complexity), we can select algorithms that are efficient, scalable, and suitable for large datasets.

1. Search Functionality

Computational Requirements

Search functionality retrieves a word, phrase, or sequence from a large collection of documents.

Requirements

- Fast response time
- Efficient pattern matching
- Ability to process millions of documents
- Low memory overhead

Alternative Algorithms

Algorithm	Time Complexity	Space Complexity	Evaluation
Linear Search	$O(n)$	$O(1)$	Simple but inefficient for large datasets
Binary Search (sorted data)	$O(\log n)$	$O(1)$	Very efficient for sorted data
KMP String Matching	$O(n + m)$	$O(m)$	Efficient for repeated pattern searches
Rabin-Karp	Average $O(n + m)$	$O(1)$	Good for multiple-pattern matching using hashing
Boyer-Moore	Best: Sub-linear, Worst: $O(nm)$	$O(\text{Alphabet Size})$	Excellent for long patterns

- Linear Search is suitable only for small datasets.
- Binary Search is effective only when data is sorted.
- KMP and Boyer-Moore are preferred for large text collections because they reduce unnecessary comparisons.

Best Choice: KMP or Boyer-Moore for large-scale text search.

2. Document Similarity Analysis

Computational Requirements

Document similarity compares two or more documents to determine how closely they are related.

Requirements

- Compare large document collections
- Efficient vector or sequence comparison
- Support semantic similarity

Alternative Algorithms

Algorithm	Time Complexity	Space Complexity	Evaluation
Brute Force Comparison	$O(n^2)$	$O(1)$	Not scalable
Edit Distance (DP)	$O(m \times n)$	$O(m \times n)$	Accurate for sequence comparison
Longest Common Subsequence (DP)	$O(m \times n)$	$O(m \times n)$	Useful for text similarity
TF-IDF + Cosine Similarity	$O(d)$	$O(d)$	Efficient for large document collections

- Brute-force comparison becomes impractical for large document collections.
- Dynamic Programming methods provide accurate similarity measurements but require more memory.
- TF-IDF with Cosine Similarity scales well for large datasets and is widely used in search engines.

Best Choice: TF-IDF + Cosine Similarity for large-scale document collections.

3. Citation Analysis

Computational Requirements

Citation analysis studies relationships among research papers to identify influential publications and knowledge flow.

Requirements

- Graph traversal
- Flow optimization
- Processing large citation networks

Brute-force comparison becomes impractical for large document collections. Dynamic Programming methods provide accurate similarity measurements but require more memory. TF-IDF with Cosine Similarity scales well for large datasets and is widely used in search engines.

Alternative Algorithms

Algorithm	Time Complexity	Space Complexity	Evaluation
DFS/BFS	$O(V + E)$	$O(V)$	Efficient for graph traversal
Ford-Fulkerson	$O(E \times \text{Max Flow})$	$O(V)$	Suitable for flow analysis
Edmonds-Karp	$O(VE^2)$	$O(V)$	More efficient than basic Ford-Fulkerson for many cases

(Here, V = number of vertices (papers), E = number of citation links.)

- BFS and DFS efficiently traverse citation graphs.
- Network Flow algorithms are preferred when analyzing knowledge flow or identifying bottlenecks.

Best Choice: Network Flow Algorithms for citation-flow analysis.

4. Scheduling Functionality

Computational Requirements

Scheduling assigns tasks or jobs to resources while minimizing completion time or maximizing resource utilization.

Requirements

- Efficient resource allocation
- Minimum waiting time
- Handling thousands of tasks

Alternative Algorithms

Algorithm	Time Complexity	Evaluation
FCFS	$O(n)$	Simple but may produce long waiting times
Shortest Job First (SJF)	$O(n \log n)$	Improves average waiting time
Priority Scheduling	$O(n \log n)$	Supports priority-based execution
Round Robin	$O(n)$	Fair allocation for time-sharing systems
Approximation Algorithms	Polynomial Time	Suitable for large NP-hard scheduling problems

- Exact scheduling becomes difficult for very large optimization problems.
 - Approximation algorithms provide near-optimal schedules within practical execution times.
- Best Choice: Approximation Algorithms** for large-scale scheduling problems.

5. Primality Testing

Computational Requirements

Primality testing determines whether a large integer is prime or composite.

Requirements

- Process numbers with hundreds or thousands of digits
- High accuracy
- Fast execution

Alternative Algorithms

Algorithm	Time Complexity	Evaluation
Trial Division	$O(\sqrt{n})$	Too slow for large numbers
Fermat Test	Polynomial (Probabilistic)	Fast but may fail for some composite numbers
Miller-Rabin	Polynomial (Probabilistic)	Very fast and widely used in cryptography
AKS Primality Test	Polynomial (Deterministic)	Guaranteed correct but slower in practice

- Trial Division is practical only for small integers.
- Miller-Rabin offers an excellent balance of speed and reliability for very large numbers.
- AKS guarantees correctness but is rarely used in real-world cryptographic systems because of its higher practical cost.

Best Choice: Miller-Rabin Primality Test.

Thank You